

Bruno Gonçalves
www.bgoncalves.com

Twitter API and
Spatial Analysis



**DATA
SCIENCE**

Anatomy of a Tweet



Anatomy of a Tweet



```
[u'contributors',
 u'truncated',
 u'text',
 u'in_reply_to_status_id',
 u'id',
 u'favorite_count',
 u'source',
 u'retweeted',
 u'coordinates',
 u'entities',
 u'in_reply_to_screen_name',
 u'in_reply_to_user_id',
 u'retweet_count',
 u'id_str',
 u'favorited',
 u'user',
 u'geo',
 u'in_reply_to_user_id_str',
 u'possibly_sensitive',
 u'lang',
 u'created_at',
 u'in_reply_to_status_id_str',
 u'place',
 u'metadata']
```


Anatomy of a Tweet

```
[u'contributors',
 u'truncated',
 u'text',
 u'in_reply_to_status_id',
 u'id',
 u'favorite_count',
 u'source',
 u'retweeted',
 u'coordinates',
 u'entities',
 u'in_reply_to_screen_name',
 u'in_reply_to_user_id',
 u'retweet_count',
 u'id_str',
 u'favorited',
 u'user',
 u'geo',
 u'in_reply_to_user_id_str',
 u'possibly_sensitive',
 u'lang',
 u'created_at',
 u'in_reply_to_status_id_str',
 u'place',
 u'metadata']

[u'follow_request_sent',
 u'profile_use_background_image',
 u'default_profile_image',
 u'id',
 u'profile_background_image_url_https',
 u'verified',
 u'profile_text_color',
 u'profile_image_url_https',
 u'profile_sidebar_fill_color',
 u'entities',
 u'followers_count',
 u'profile_sidebar_border_color',
 u'id_str',
 u'profile_background_color',
 u'listed_count',
 u'is_translation_enabled',
 u'utc_offset',
 u'statuses_count',
 u'description',
 u'friends_count',
 u'location',
 u'profile_link_color',
 u'profile_image_url',
 u'following',
 u'geo_enabled',
 u'profile_banner_url',
 u'profile_background_image_url',
 u'screen_name',
 u'lang',
 u'profile_background_tile',
 u'favourites_count',
 u'name',
 u'notifications',
 u'url',
 u'created_at',
 u'contributors_enabled',
 u'time_zone',
 u'protected',
 u'default_profile',
 u'is_translator']
```

Anatomy of a Tweet

```
[u'contributors',  
 u'truncated',  
 u'text',  
 u'in_reply_to_status_id',  
 u'id',  
 u'favorite_count',  
 u'source',  
 u'retweeted',  
 u'coordinates',  
 u'entities',  
 u'in_reply_to_screen_name',  
 u'in_reply_to_user_id',  
 u'retweet_count',  
 u'id_str',  
 u'favorited',  
 u'user',  
 u'geo',  
 u'in_reply_to_user_id_str',  
 u'possibly_sensitive',  
 u'lang',  
 u'created_at',  
 u'in_reply_to_status_id_str',  
 u'place',  
 u'metadata']
```

```
u"I'm at Terminal Rodovi\xelrio de Feira de Santana  
 (Feira de Santana, BA) http://t.co/WirvdHwYMq"
```

```
u"<a href='\"http://foursquare.com\"' rel='\"nofollow\"'  
 foursquare</a>"
```

```
[u'symbols',  
 u'user_mentions',  
 u'hashtags',  
 u'urls']
```

```
[u'type',  
 u'coordinates']
```

Anatomy of a Tweet


```
[u'contributors',
 u'truncated',
 u'text',
 u'in_reply_to_status_id',
 u'id',
 u'favorite_count',
 u'source',
 u'retweeted',
 u'coordinates',
 u'entities',
 u'in_reply_to_screen_name',
 u'in_reply_to_user_id',
 u'retweet_count',
 u'id_str',
 u'favorited',
 u'user',
 u'geo',
 u'in_reply_to_user_id_str',
 u'possibly_sensitive',
 u'lang',
 u'created_at',
 u'in_reply_to_status_id_str',
 u'place',
 u'metadata']
```

```
u"I'm at Terminal Rodoviário de Feira de Santana
 (Feira de Santana, BA) http://t.co/WirvdHwYMq"
```

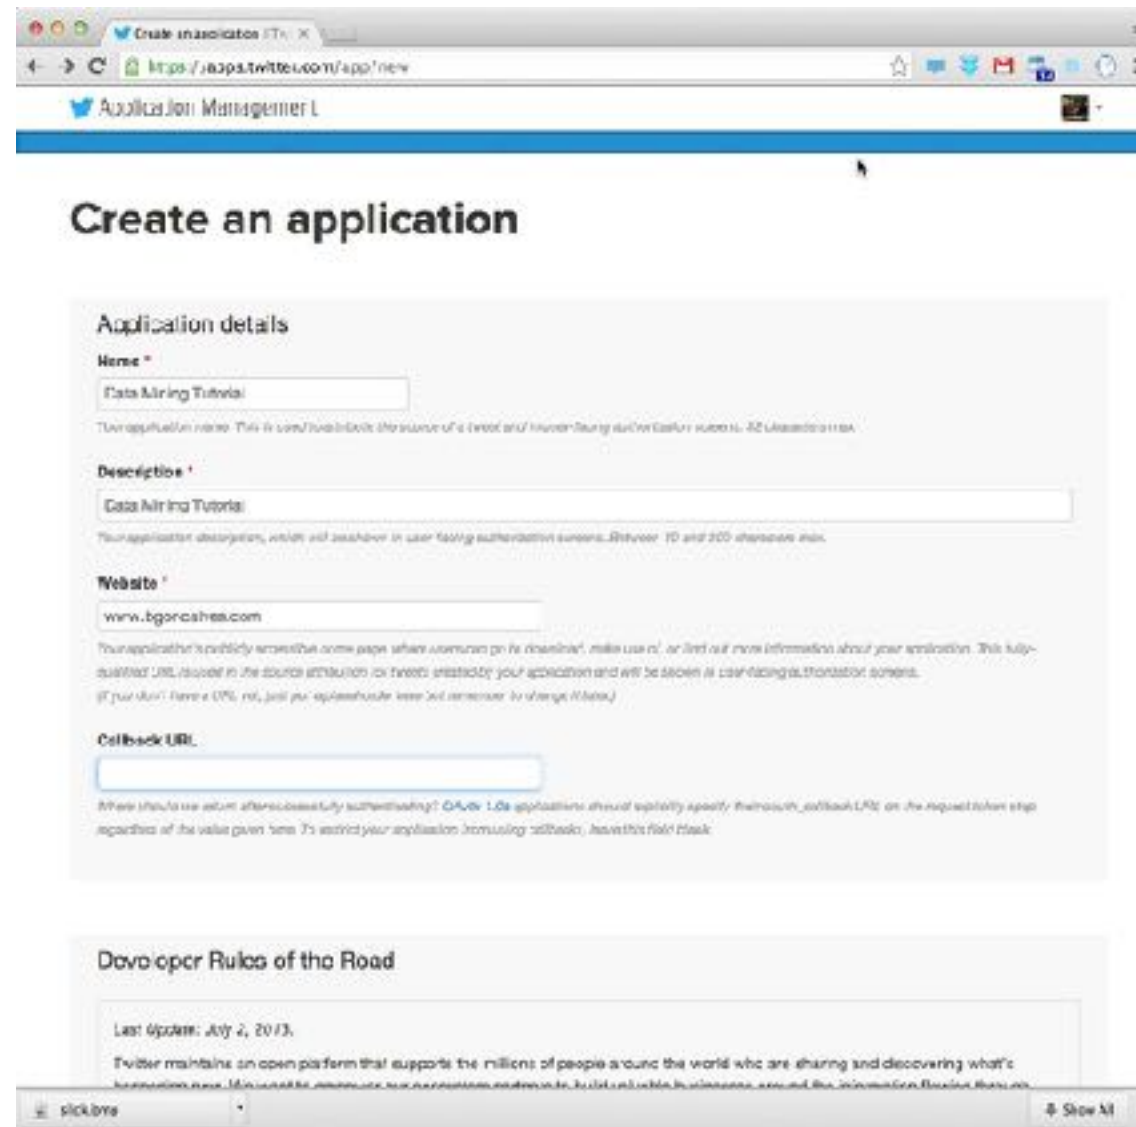
```
u'<a href='\"http://foursquare.com\" rel='\"nofollow\">
 foursquare</a>'
```

```
[u'symbols',
 u'user_mentions',
 u'hashtags',
 u'urls', {u'display_url': u'4sq.com/1k5MeYF',
 u'expanded_url': u'http://4sq.com/1k5MeYF',
 u'indices': [70, 92],
 u'url': u'http://t.co/WirvdHwYMq'}]
[u'type',
 u'coordinates']
```

Registering an Application



Registering an Application



The screenshot shows a web browser window with the URL <https://apps.twitter.com/new>. The page title is "Application Manager L". The main heading is "Create an application".

Application details

Name *

Your application name. This is used to identify the source of a tweet and is visible during OAuth requests. 32 characters max.

Description *

Your application description, which will appear in user-facing authorization screens. Between 70 and 200 characters max.

Website *

Your application's publicly accessible home page where users can go to download, make use of, or find out more information about your application. This fully-qualified URL, located in the source attribution on tweets created by your application and will be shown in user-facing authorization screens. (If you don't have a URL yet, just put a placeholder here and remember to change it later.)

Callback URL

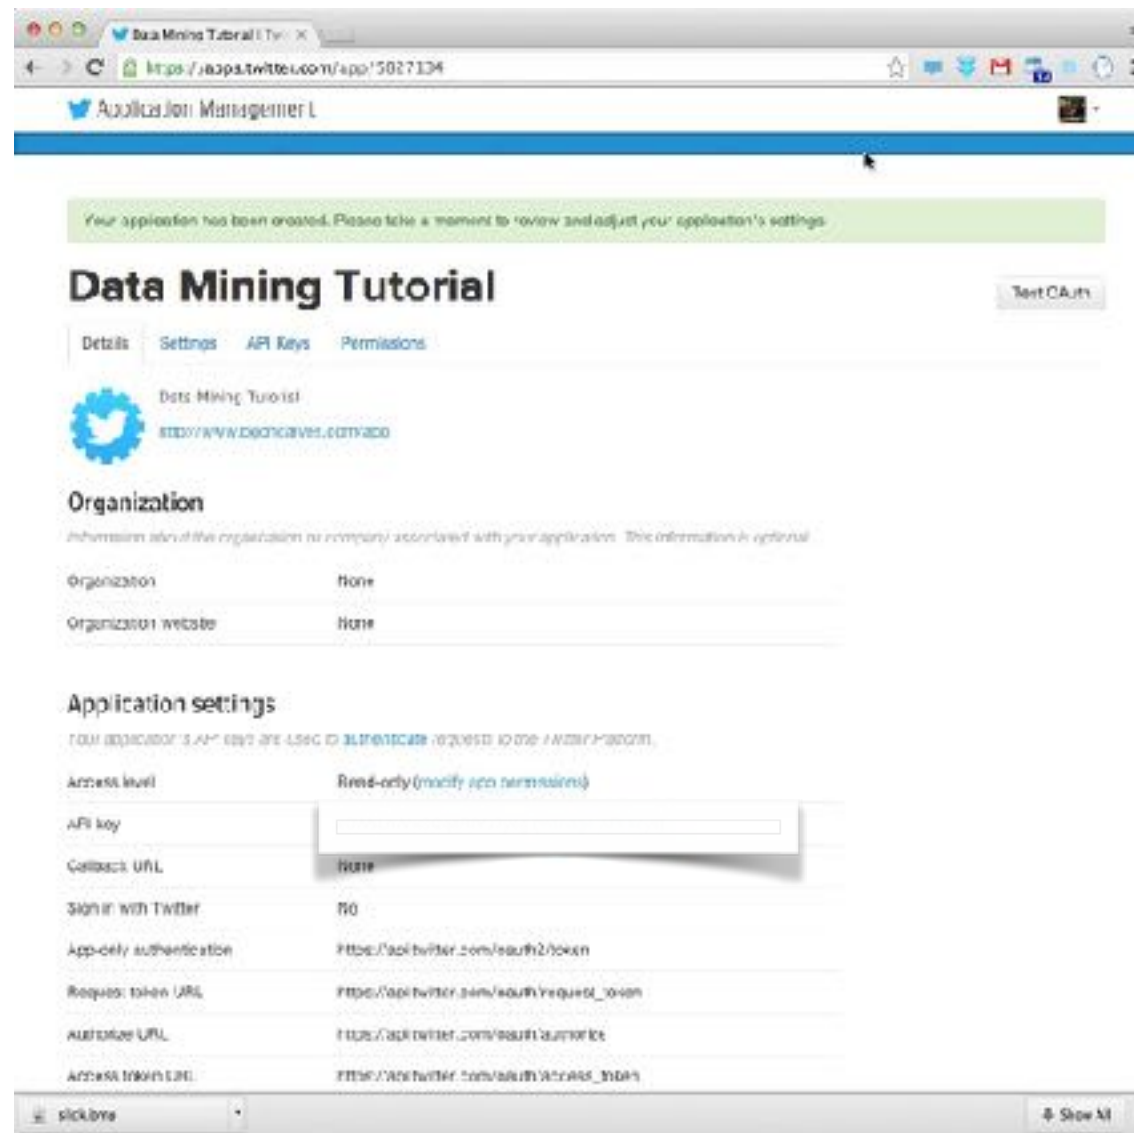
Where should we return after successfully authenticating? OAuth 1.0a applications should explicitly specify their oauth_callback URL on the request token step regardless of the value given here. To restrict your application (preventing callbacks), leave this field blank.

Developer Rules of the Road

Last Update: July 2, 2013.
Twitter maintains an open platform that supports the millions of people around the world who are sharing and discovering what's interesting, new, fun and useful. To ensure the platform remains open, we've established a set of rules that govern how developers can interact with the platform. These rules are designed to protect the integrity of the platform and ensure that all developers are treated fairly.

pick.me Show All

Registering an Application



The screenshot shows the Twitter Application Management interface for an application named 'Data Mining Tutorial'. The page is titled 'Application Management' and includes a green notification bar stating 'Your application has been created. Please take a moment to review and adjust your application's settings.' Below this, the application name 'Data Mining Tutorial' is displayed with a 'Test OAuth' button. The 'Details' tab is selected, showing the application's profile picture, name, and website. The 'Organization' section is currently empty. The 'Application settings' section contains a table of configuration options.

Application settings	
YOUR APPLICATION'S API KEY IS USED TO authenticate requests to the / API endpoint.	
Access level	Read-only (modify app permissions)
API key	
Contact URL	None
Sign in with Twitter	No
App-only authentication	https://api.twitter.com/oauth2/token
Request token URL	https://api.twitter.com/oauth/request_token
Authorize URL	https://api.twitter.com/oauth/authorize
Access token URL	https://api.twitter.com/oauth/access_token

Registering an Application

The screenshot shows a web browser window with the URL `https://apps.twitter.com/app/5027134/keys`. The page title is "Data Mining Tutorial" and it includes a "Test OAuth" button. Below the title are tabs for "Details", "Settings", "API Keys", and "Permissions", with "Settings" being the active tab. The "Application settings" section contains a warning about the API secret, followed by fields for "API key" (which is highlighted with a red box), "API secret", "Access level" (set to "Read-only (no write access permissions)", "Owner" (set to "bgoncalves"), and "Owner ID" (set to "15308566"). Below this is the "Application actions" section with buttons for "Regenerate API keys" and "Change App Permissions". The "Your access token" section includes a warning about authorizing the application and a "Create my access token" button. At the bottom, there is a "Token actions" section with a "Create my access token" button. The browser's address bar and tabs are visible at the top, and a "Show All" button is in the bottom right corner.

Data Mining Tutorial [Test OAuth](#)

[Details](#) [Settings](#) [API Keys](#) [Permissions](#)

Application settings

Keep this 'API secret' a secret. This key should never be human-readable to your application.

API key	<input type="text"/>
API secret	<input type="text"/>
Access level	Read-only (no write access permissions)
Owner	bgoncalves
Owner ID	15308566

Application actions

[Regenerate API keys](#) [Change App Permissions](#)

Your access token

You haven't authorized this application for your own account yet.

By creating your access token link, you will have everything you need to make API calls right away. The access token generated will be assigned your application's current permission level.

Token actions

[Create my access token](#)

[Show All](#)

Registering an Application

The screenshot shows the 'Data Mining Tutorial' application settings page in the Twitter Developer Portal. The page is titled 'Data Mining Tutorial' and has a 'Test OAuth' button. Below the title are tabs for 'Details', 'Settings', 'API Keys', and 'Permissions'. The 'API Keys' tab is selected, showing 'Application settings'. The settings include: 'API key' (a text input field), 'API secret' (a text input field), 'Access level' (a dropdown menu set to 'Read-only (no write API permissions)'), 'Owner' (a dropdown menu set to 'Egonwales'), and 'Owner ID' (a text input field containing '15068566'). Below these settings are two buttons: 'Regenerate API keys' and 'Change App Permissions'. The 'Your access token' section is also visible, with a note: 'This access token can be used to make API requests on your own account's behalf. Do not share your access token secret with anyone.' It includes fields for 'Access token' (a text input field), 'Access token secret' (a text input field), 'Access level' (a dropdown menu set to 'Read-only'), 'Owner' (a dropdown menu set to 'Egonwales'), and 'Owner ID' (a text input field containing '15068566'). The page is viewed in a browser window with the address bar showing 'https://apps.twitter.com/api/1.1/keys'.

Data Mining Tutorial Test OAuth

[Details](#) [Settings](#) [API Keys](#) [Permissions](#)

Application settings

Keep this "API secret" a secret. This key should never be human-readable in your application.

API key

API secret

Access level Read-only (no write API permissions)

Owner Egonwales

Owner ID 15068566

Application actions

[Regenerate API keys](#) [Change App Permissions](#)

Your access token

This access token can be used to make API requests on your own account's behalf. Do not share your access token secret with anyone.

Access token

Access token secret

Access level Read-only

Owner Egonwales

Owner ID 15068566

[Show All](#)

API Basics

<https://dev.twitter.com/docs>

- The **twitter** module provides the oauth interface. We just need to provide the right credentials.
- Best to keep the credentials in a **dict** and parametrize our calls with the dict key. This way we can switch between different accounts easily.
- **.Twitter(auth)** takes an **OAuth** instance as argument and returns a **Twitter** object that we can use to interact with the API
- **Twitter** methods mimic API structure
- 4 basic types of objects:
 - Tweets
 - Users
 - Entities

Authenticating with the API

```
import tweepy
from twitter_accounts import accounts

app = accounts["social"]

auth = twitter.oauth.OAuth(app["token"],
                             app["token_secret"],
                             app["api_key"],
                             app["api_secret"])

twitter_api = twitter.Twitter(auth=auth)
```

- In the remainder of this course, the **accounts** dict will live inside the `twitter_accounts.py` file
- 4 basic types of objects:
 - Tweets
 - Users
 - Entities
 - Places

Searching for Tweets

<https://dev.twitter.com/docs/api/1.1/get/search/tweets>

- `.search.tweets(query, count)`
 - `query` is the content to search for
 - `count` is the maximum number of results to return
- returns dict with a list of “`statuses`” and “`search_metadata`”

```
{u'completed_in': 0.027,  
 u'count': 15,  
 u'max_id': 438088492577345536,  
 u'max_id_str': u'438088492577345536',  
 u'next_results': u'?max_id=438088485145034752&q=soccer&include_entities=1',  
 u'query': u'soccer',  
 u'refresh_url': u'?since_id=438088492577345536&q=soccer&include_entities=1',  
 u'since_id': 0,  
 u'since_id_str': u'0'}
```
- `search_results[“search_metadata”][“next_results”]` can be used to get the next page of results

Streaming Geocoded data

<https://dev.twitter.com/streaming/overview/request-parameters#locations>

- The Streaming api provides realtime data, subject to filters
- Use `TwitterStream` instead of `Twitter` object (`.TwitterStream(auth=twitter_api.auth)`)
- `.status.filter(track=q)` will return tweets that match the query `q` in real time
- Returns generator that you can iterate over
- `.status.filter(locations=bb)` will return tweets that occur within the bounding box `bb` in real time
- `bb` is a comma separated pair of lon/lat coordinates.
 - -180,-90,180,90 - World
 - -74,40,-73,41 - NYC

Streaming Geocoded data

<https://dev.twitter.com/streaming/overview/request-parameters#locations>

```
import twitter
from twitter_accounts import accounts
import sys
import gzip

app = accounts["social"]

auth = twitter.oauth.OAuth(app["token"],
                             app["token_secret"],
                             app["api_key"],
                             app["api_secret"])

stream_api = twitter.TwitterStream(auth=auth)

query = "-74,40,-73,41" # NYC
stream_results = stream_api.statuses.filter(locations=query)
tweet_count = 0

fp = gzip.open("NYC.json.gz", "a")

for tweet in stream_results:
    try:
        tweet_count += 1
        print(tweet_count, tweet["id"])
        print(tweet, file=fp)
    except:
        pass

if tweet_count % 10000 == 0:
    print(tweet_count, file=sys.stderr)
    break
```

Plotting geolocated tweets

```
import sys
import gzip
import matplotlib.pyplot as plt

x = []
y = []

line_count = 0


try:
    for line in gzip.open(sys.argv[1]):
        try:
            tweet = eval(line.strip())
            line_count += 1

            if "coordinates" in tweet and tweet["coordinates"] is not None:
                x.append(tweet["coordinates"]["coordinates"][0])
                y.append(tweet["coordinates"]["coordinates"][1])
        except:
            pass
    except:
        pass

print("Read", line_count, "and found", len(x), "geolocated tweets", file=sys.stderr)

plt.plot(x, y, '*')
plt.xlabel('Longitude')
plt.ylabel('Latitude')
plt.savefig(sys.argv[1] + '.png')
plt.close()
```

Plotting geolocated tweets



Shapefiles

<http://www.esri.com/library/whitepapers/pdfs/shapefile.pdf>

- Open specification developed by ESRI, still the current leader in commercial GIS software
- shapefiles aren't actual (individual) files...
- but actually a set of files sharing the same name but with different extensions:

```
(py35) (master) bgoncalves@underdark:$ls -l
total 4856
-rw-r--r--@ 1 bgoncalves  staff      537 Apr 17 12:40 nybb.dbf
-rw-r--r--@ 1 bgoncalves  staff      562 Apr 17 12:40 nybb.prj
-rw-r--r--@ 1 bgoncalves  staff 1217376 Apr 17 12:40 nybb.shp
-rw-r--r--@ 1 bgoncalves  staff    12905 Apr 17 12:40 nybb.shp.xml
-rw-r--r--@ 1 bgoncalves  staff     140 Apr 17 12:40 nybb.shx
-rw-r--r--  1 bgoncalves  staff     536 Apr 17 12:40 nybb_wgs84.dbf
-rw-r--r--  1 bgoncalves  staff     143 Apr 17 12:40 nybb_wgs84.prj
-rw-r--r--  1 bgoncalves  staff     257 Apr 17 12:40 nybb_wgs84.qpj
-rw-r--r--  1 bgoncalves  staff 1217376 Apr 17 12:40 nybb_wgs84.shp
-rw-r--r--  1 bgoncalves  staff     140 Apr 17 12:40 nybb_wgs84.shx
(py35) (master) bgoncalves@underdark:$
```


- the actual set of files changes depending on the contents, but three files are usually present:
 - **.shp** - also commonly referred to as “the” shapefile. Contains the geometric information
 - **.dbf** - a simple database containing the feature attribute table.
 - **.shx** - a spatial index, not strictly required

Shapefiles

http://www.nyc.gov/html/dcp/html/bytes/districts_download_metadata.shtml#bcd

Borough Boundaries & Community Districts	Download	Metadata
Borough Boundaries (Clipped to Shoreline)	 (645k)	
Borough Boundaries (Water Areas Included)	 (31k)	
Community Districts (Clipped to Shoreline)	 (772k)	

- Unfortunately it doesn't use the right reference system (**WGS84**), so we must convert it.



pyshp

<https://github.com/GeospatialPython/pyshp>

- **pyshp** defines utility functions to load and manipulate Shapefiles programmatically.
- The **shapefile** module handles the most common operations:
 - **.Reader(filename)** - Returns a **Reader** object
- **Reader.records()/Reader.iterRecords()** returns/iterates over the different records present in the shapefile
- **Reader.shapes()/Reader.iterShapes()** - returns/iterates over the different shapes present in the shapefile
- **Reader.shapeRecords()/Reader.iterShapeRecords()** returns/iterates over both shapes and records present in the shapefile
- **Reader.record(index)/Reader.shape(index)/Reader.shapeRecord(index)** - return the record/shape/shapeRecord at index position **index**
- **Reader.numRecords** - returns the number of records in the shapefile

```
import sys
import shapefile

shp = shapefile.Reader('geofiles/nybb_15c/nybb_wgs84.shp')

print("Found", shp.numRecords, "records:")

recordDict = dict(zip([record[1] for record in shp.iterRecords()], range(shp.numRecords)))

for record, id in recordDict.items():
    print(id, record)
```

- **shape** objects contain several fields:
 - **bbox** - lower left and upper right **x,y** coordinates (long/lat) - **optional**
 - **parts** - list of indexes for the first point of each of the parts making up the shape.
 - **points** - **x,y** coordinates for each point in the shape.
 - **shapeType** - integer representing the shape type - all shapes in a shapefile are required to be of the same **shapeType** or **null**.

Value	Shape Type
0	Null Shape
1	Point
3	PolyLine
5	Polygon
8	MultiPoint
11	PointZ
13	PolyLineZ
15	PolygonZ
18	MultiPointZ
21	PointM
23	PolyLineM
25	PolygonM
28	MultiPointM
31	MultiPatch

Challenge - pyshp

- Write a simple script to plot out all the shapes in:

geofiles/nybb_15c/nybb_wgs84.shp

```
import shapefile
import matplotlib.pyplot as plt
import numpy as np


shp = shapefile.Reader('geofiles/nybb_15c/nybb_wgs84.shp')

pos = None
count = 0
for shape in shp.iterShapes():
    points = np.array(shape.points)
    parts = shape.parts
    parts.append(len(shape.points))

    for i in range(len(parts)-1):
        plt.plot(points.T[0][parts[i]:parts[i+1]], points.T[1][parts[i]:parts[i+1]])

plt.savefig('NYC.png')
```

Challenge - pyshp



shapely

<http://toblerity.org/shapely/manual.html>

- Shapely defines geometric objects under `shapely.geometry`:
 - Point
 - Polygon
 - MultiPolygon
- `shape()` Convenience function that creates the appropriate geometric object
- and common operations
 - `.crosses(shape)` - if it partially overlaps `shape`
 - `.contains(shape)` - whether it contains or not the object `shape`
 - `.within(shape)` - whether it is contained by object `shape`
 - `.touches(shape)` - if the boundaries of this object touch `shape`

shapely

<http://toblerity.org/shapely/manual.html>

- **shape** objects provide useful fields to query a shapes properties:
 - **.centroid** - The centroid ("center of mass") of the object
 - **.area** - returns the area of the object
 - **.bounds** - the MBR of the shape in (minx, miny, maxx, maxy) format
 - **.length** - the length of the shape
 - **.geom_type** - the Geometry Type of the object
- **shapely.shape** is also able to easily load **pyshp**'s shape objects to allow for further manipulations.

```
import sys
import shapefile
from shapely.geometry import shape

shp = shapefile.Reader('geofiles/nybb_15c/nybb_wgs84.shp')

recordDict = dict(zip([record[1] for record in shp.iterRecords()], range(shp.numRecords)))

manhattan = shape(shp.shape(recordDict["Manhattan"]))

print("Centroid:", manhattan.centroid)
print("Bounding box:", manhattan.bounds)
print("Geometry type:", manhattan.geom_type)
print("Length:", manhattan.length)
```


Challenge - Filter points within a Shapefile

- Load each tweet in NYC.json.gz file by using:

```
tweet = eval(line.strip())

import sys
import shapefile
• from shapely.geometry import shape, Point
import gzip

shp = shapefile.Reader('geofiles/nybb_15c/nybb_wgs84.shp')

recordDict = dict(zip([record[1] for record in shp.iterRecords()], range(shp.numRecords)))

manhattan = shape(shp.shape(recordDict["Manhattan"]))
fp = gzip.open("Manhattan.json.gz", "w")


for line in gzip.open("NYC.json.gz"):
    try:
        tweet = eval(line.strip())

        if "coordinates" in tweet and tweet["coordinates"] is not None:
            point = Point(tweet["coordinates"]["coordinates"])

            if manhattan.contains(point):
                fp.write(line)
    except:
        pass

fp.close()
```

Challenge - Filter points within a Shapefile





Twitter places

- As we saw last week, Twitter defines a “**coordinates**” field in tweets
- There is also a “**place**” field that we glossed over.
- The **place** object contains also geographical information, but at a coarser resolution than the **coordinates** field.
- Each place has a unique **place_id**, a **bounding_box** and some geographical information, such as **country** and **full_name**:

```
{ 'attributes': {},  
  'bounding_box': { 'coordinates': [[[-74.041878, 40.570842],  
    [-74.041878, 40.739434],  
    [-73.855673, 40.739434],  
    [-73.855673, 40.570842]]],  
    'type': 'Polygon'},  
  'country': 'United States',  
  'country_code': 'US',  
  'full_name': 'Brooklyn, NY',  
  'id': '011add077f4d2da3',  
  'name': 'Brooklyn',  
  'place_type': 'city',  
  'url': 'https://api.twitter.com/1.1/geo/id/011add077f4d2da3.json' }
```

- places can be of several different types: **'admin'**, **'city'**, **'neighborhood'**, **'poi'**



Twitter places

- As we already saw, Twitter defines a “**coordinates**” field in tweets
- There is also a “**place**” field that we glossed over.
- The **place** object contains also geographical information, but at a coarser resolution than the **coordinates** field.
- Each place has a unique **place_id**, a **bounding_box** and some geographical information, such as **country** and **full_name**:

```
{ 'attributes': {},  
  'bounding_box': { 'coordinates': [[[-74.041878, 40.570842],  
    [-74.041878, 40.739434],  
    [-73.855673, 40.739434],  
    [-73.855673, 40.570842]]],  
    'type': 'Polygon'},  
  'country': 'United States',  
  'country_code': 'US',  
  'full_name': 'Brooklyn, NY',  
  'id': '011add077f4d2da3',  
  'name': 'Brooklyn',  
  'place_type': 'city',  
  'url': 'https://api.twitter.com/1.1/geo/id/011add077f4d2da3.json' }
```

The bounding_box field is GeoJSON formatted and compatible with `pyshp.shape`

- places can be of several different types: **'admin'**, **'city'**, **'neighborhood'**, **'poi'**

Twitter places

<https://dev.twitter.com/overview/api/places>

Place Attributes

Place Attributes are metadata about places. An attribute is a key-value pair of arbitrary strings, but with some conventions.

Below are a number of well-known place attributes which may, or may not exist in the returned data. These attributes are provided when the place was created in the Twitter places database.

Key	Description
street_address	
locality	the city the place is in
region	the administrative region the place is in
iso3	the country code
postal_code	in the preferred local format for the place
phone	in the preferred local format for the place, include long distance code
twitter	twitter screen-name, without @
url	official/canonical URL for place
app:id	An ID or comma separated list of IDs representing the place in the applications place database.

Keys can be no longer than 140 characters in length. Values are unicode strings and are restricted to 2000 characters.

Challenge - Filter points and places

- Load each tweet in **NYC.json.gz** file by using:

```
tweet = eval(line.strip())
```

- on each line and write to **Manhattan_places.json.gz** all the tweets within Manhattan, as defined by

```
geofiles/nybb_15c/nybb_wgs84.shp
```

- but now check if the centroid of the **'place'** object **'bounding_box'** is within Manhattan as well

Challenge - Filter points and places

```
import sys
import shapefile
from shapely.geometry import shape, Point
import gzip

shp = shapefile.Reader('geofiles/nybb_15c/nybb_wgs84.shp')

recordDict = dict(zip([record[1] for record in shp.iterRecords()], range(shp.numRecords)))

manhattan = shape(shp.shape(recordDict["Manhattan"]))
fp = gzip.open("Manhattan_places.json.gz", "w")


for line in gzip.open("NYC.json.gz"):
    try:
        tweet = eval(line.strip())
        point = None

        if "coordinates" in tweet and tweet["coordinates"] is not None:
            point = Point(tweet["coordinates"]["coordinates"])
        else:
            if "place" in tweet and tweet["place"]["bounding_box"] is not None:
                bbox = shape(tweet["place"]["bounding_box"])
                point = bbox.centroid

        if point is not None and manhattan.contains(point):
            fp.write(line)
    except:
        pass

fp.close()
```


Challenge - Filter points and places



Challenge - Filter points and places

```
import sys
import gzip
import numpy as np
import shapefile
from shapely.geometry import shape, Point
import matplotlib.pyplot as plt

shp = shapefile.Reader('geofiles/nybb_15c/nybb_wgs84.shp')
recordDict = dict(zip([record[1] for record in shp.iterRecords()],
range(shp.numRecords)))

manhattan = shp.shape(recordDict["Manhattan"])

points = np.array(manhattan.points)
parts = manhattan.parts
parts.append(len(manhattan.points))

for i in range(len(parts)-1):
    plt.plot(points.T[0][parts[i]:parts[i+1]], points.T[1]
[parts[i]:parts[i+1]])

points_X = []
points_Y = []

for line in gzip.open(sys.argv[1]):
    try:
        tweet = eval(line.strip())
        point = None

        if "coordinates" in tweet and tweet["coordinates"] is not None:
            point = Point(tweet["coordinates"]["coordinates"])
        else:
            if "place" in tweet and tweet["place"]["bounding_box"] is not
None:
                bbox = shape(tweet["place"]["bounding_box"])
                point = bbox.centroid

            if point is not None:
                points_X.append(point.x)
                points_Y.append(point.y)
    except:
        pass

plt.plot(points_X, points_Y, 'r+')
plt.savefig(sys.argv[1] + '.png')
```

Calculating distances

https://en.wikipedia.org/wiki/Vincenty%27s_formulae
https://en.wikipedia.org/wiki/Great-circle_distance
https://en.wikipedia.org/wiki/Haversine_formula

- Earlier we saw how to obtain the distance between two points using the Google Maps API.
- But what is the shortest distance between two **arbitrary** points on the surface of the Earth?
- This depends strongly on our model of the Earth:
 - **Great Circle** - Assumes that the Earth is a perfect sphere of a given radius
 - Usually uses the **Haversine** formula
$$\Delta\sigma = 2 \arcsin \sqrt{\sin^2\left(\frac{\Delta\phi}{2}\right) + \cos\phi_1 \cdot \cos\phi_2 \cdot \sin^2\left(\frac{\Delta\lambda}{2}\right)}$$
 - **Vincenty** - Uses a (more) accurate ellipsoid model of the Earth

geopy

<https://geopy.readthedocs.io/en/1.10.0/>

- `geopy` provides two different types of functionality
 - `geopy.geocoders` - a unified interface to several geocoding services (Google Maps, Nominatim, Yahoo, Bing, etc...)
 - `geopy.distance` - state of the art distance calculations
- We will focus just on the `distance` module:
 - `distance.vincenty(p1, p2)` - Calculate the vincenty distance between `p1` and `p2`
 - `distance.great_circle(p1, p2)` - Calculate the great circle distance between `p1` and `p2`
 - `distance.distance(p1, p2)` - an alias to `distance.vincenty` to be used as a default.

- all `distance` functions return a `Distance` object.
- the `Distance` object provides properties that represent the result in different units:
 - `.km/.kilometers`
 - `.m/.meters`
 - `.mi/.miles`
 - `.ft/.feet`
 - `.nm/.nautical`
- it also allows us to recalculate the result using different ellipsoids:
 - `.set_ellipsoid('ellipsoid')`
 - by default `WGS-84` is used.

```
ELLIPSOIDS = {
    'model':
    'major (km)':
    'minor (km)':
    'flattening':
    'WGS-84': (6378.137, 6356.7523142, 1 /
    'GRS-80': (6378.137, 6356.7523141, 1 /
    'Airy (1830)': (6377.563396, 6356.256909, 1 /
    'Intl 1924': (6378.388, 6356.911946, 1 / 297.0),
    'Clarke (1880)': (6378.249145, 6356.51486955, 1 / 293.465),
    'GRS-67': (6378.1600, 6356.774719, 1 / 298.25),
}
```

Challenge - geopy

<https://geopy.readthedocs.io/en/1.10.0/>

- Calculate the distance between

p1 = (41.49008, -71.312796)
p2 = (41.499498, -81.695391)

- in meters, using the `vincenty` and `great_circle` functions.

```
from geopy import distance

p1 = (41.49008, -71.312796)
p2 = (41.499498, -81.695391)

dist_vincenty = distance.vincenty(p1,
p2).meters
dist_great = distance.great_circle(p1,
p2).meters

print("Vincenty:", dist_vincenty)
print("Great Circles:", dist_great)
```

ASCII Grid

- Perhaps the simplest raster file
- ASCII text based
- A small header

```
ncols      246
nrows      119
xllcorner  -126.500000000000
yllcorner   22.750000000000
cellsize    0.250000000000
NODATA_value -9999
```

- Followed by rows of numbers
- Very convenient to Read and Write

Challenge - ASCII Grid

- Write a simple function to load the file:

`geofiles/US_pop.asc`

- into a numpy array and print the value corresponding to:

`40.730503, -74.243251`

Challenge

```
import numpy as np
import matplotlib.pyplot as plt

def map_points(xllcorner, yllcorner, cellsize, nrows, x, y):
    x = int((x-xllcorner)/cellsize)
    y = (nrows-1)-int((y-yllcorner)/cellsize)

    return x, y

fp = open("geofiles/US_pop.asc")
ncols, count = fp.readline().split()
ncols = int(count)
nrows, count = fp.readline().split()
nrows = int(count)
xllcorner, value = fp.readline().split()
xllcorner = float(value)
yllcorner, value = fp.readline().split()
yllcorner = float(value)
cellsize, value = fp.readline().split()
cellsize = float(value)

NODATA_value, value = fp.readline().split()
NODATA_value = float(value)

data = []
for line in fp:
    fields = line.strip().split()
    data.append([float(field) for field in fields])

data = np.array(data)
data[data==NODATA_value] = 0

x = -74.243251
y = 40.730503

coord_x, coord_y = map_points(xllcorner, yllcorner, cellsize, nrows, x, y)
print(data[coord_y, coord_x])
```

Challenge

```
import numpy as np
import matplotlib.pyplot as plt

def map_points(xllcorner, yllcorner, cellsize, nrows, x, y):
    x = int((x-xllcorner)/cellsize)
    y = (nrows-1)-int((y-yllcorner)/cellsize)

    return x, y

fp = open("geofiles/US_pop.asc")
ncols, count = fp.readline().split()
ncols = int(count)
nrows, count = fp.readline().split()
nrows = int(count)
xllcorner, value = fp.readline().split()
xllcorner = float(value)
yllcorner, value = fp.readline().split()
yllcorner = float(value)
cellsize, value = fp.readline().split()
cellsize = float(value)

NODATA_value, value = fp.readline().split()
NODATA_value = float(value)


data = []
for line in fp:
    fields = line.strip().split()
    data.append([float(field) for field in fields])

data = np.array(data)
data[data==NODATA_value] = 0

x = -74.243251
y = 40.730503

coord_x, coord_y = map_points(xllcorner, yllcorner, cellsize, nrows, x, y)
print(data[coord_y, coord_x])
```

Challenge - ASCII Grid



ASCII Grid

- This type of grid is a very convenient way to aggregate spatial data.
- Simply map **lat, lon** pairs to matrix entries and then increment the values
- All we need is to define the **bbox** we are interested in, and the size of each cell and create a matrix with that shape.

```
import numpy as np
from shapely.geometry import shape, Point
import shapefile

shp = shapefile.Reader('../Lecture IV/geofiles/nybb_15c/nybb_wgs84.shp')
recordDict = dict(zip([record[1] for record in shp.iterRecords()],
                      range(shp.numRecords)))

manhattan = shp.shape(recordDict["Manhattan"])

xllcorner, yllcorner, xurcorner, yurcorner = manhattan.bbox

cellsize = 0.01

ncols = int((xurcorner-xllcorner)/cellsize)
nrows = int((yurcorner-yllcorner)/cellsize)

data = np.zeros((nrows, ncols), dtype='int')

print(data.shape)
```

Challenge - Aggregate

- Rewrite

`shapefile_filter_places.py`

- so that it now counts how many tweets happen in each cell of width **0.01**

Challenge - Aggregate

```
import sys
import numpy as np
import shapefile
from shapely.geometry import shape, Point
import matplotlib.pyplot as plt
import gzip

def map_points(xllcorner, yllcorner, cellsize, nrows, x, y):
    x = int((x-xllcorner)/cellsize)
    y = (nrows-1)-int((y-yllcorner)/cellsize)

    return x, y

def save_asc(data, xllcorner, yllcorner, cellsize, filename):
    fp = open(filename, "w")

    nrows, ncols = data.shape

    print("ncols", ncols, file=fp)
    print("nrows", nrows, file=fp)
    print("xllcorner", xllcorner, file=fp)
    print("yllcorner", yllcorner, file=fp)
    print("cellsize", cellsize, file=fp)
    print("NODATA_value", "-9999", file=fp)

    for i in range(nrows):
        for j in range(ncols):
            print((" %u " % data[i, j]), end="", file=fp)

        print("\n", end="", file=fp)

    fp.close()
```

```

shp = shapefile.Reader('../Lecture IV/geofiles/nybb_15c/nybb_wgs84.shp')
recordDict = dict(zip([record[1] for record in shp.iterRecords()], range(shp.numRecords)))
manhattan = shape(shp.shape(recordDict["Manhattan"]))

xllcorner, yllcorner, xurcorner, yurcorner = manhattan.bounds
cellsize = 0.01

ncols = int((xurcorner-xllcorner)/cellsize)
nrows = int((yurcorner-yllcorner)/cellsize)

data = np.zeros((nrows, ncols), dtype='int')

for line in gzip.open("../Lecture IV/NYC.json.gz"):
    try:
        tweet = eval(line.strip())
        point = None

        if "coordinates" in tweet and tweet["coordinates"] is not None:
            point = Point(tweet["coordinates"]["coordinates"])
        else:
            if "place" in tweet and tweet["place"]["bounding_box"] is not None:
                bbox = shape(tweet["place"]["bounding_box"])
                point = bbox.centroid

        if point is not None and manhattan.contains(point):
            coord_x, coord_y = map_points(xllcorner, yllcorner, cellsize, nrows, point.x, point.y)
            data[coord_y, coord_x] += 1

    except:
        pass

save_asc(data, xllcorner, yllcorner, cellsize, "Manhattan.asc")

plt.imshow(np.log(data+1))
plt.colorbar()
plt.savefig('Manhattan_cells.png')

```


Challenge - Aggregate

